

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
data=pd.read_csv('spam.csv')
data.head()

```

	Unnamed: 0	label	text
0	605	ham	Subject: enron methanol ; meter # : 988291\r\n...
1	2349	ham	Subject: hpl nom for january 9 , 2001\r\n( see...
2	3624	ham	Subject: neon retreat\r\nho ho ho , we ' re ar...
3	4685	spam	Subject: photoshop , windows , office . cheap ...
4	2030	ham	Subject: re : indian springs\r\nthis deal is t...

```

label_num
0      0
1      0
2      0
3      1
4      0

```

```
print(len(data))
```

```
5171
```

```
s=data['label'].value_counts()
```

```
data['label_num'].value_counts()
```

```
0      3672
```

```
1      1499
```

```
Name: label_num, dtype: int64
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 5171 entries, 0 to 5170
```

```
Data columns (total 4 columns):
```

#	Column	Non-Null Count	Dtype
0	Unnamed: 0	5171 non-null	int64
1	label	5171 non-null	object
2	text	5171 non-null	object

```
3  label_num  5171 non-null  int64
dtypes: int64(2), object(2)
memory usage: 161.7+ KB
```

```
data.describe
```


```
<bound method NDFrame.describe of      Unnamed: 0 label
text \
0          605   ham  Subject: enron methanol ; meter # : 988291\r\
n...
1          2349   ham  Subject: hpl nom for january 9 , 2001\r\
n( see...
2          3624   ham  Subject: neon retreat\r\nho ho ho , we ' re
ar...
3          4685  spam  Subject: photoshop , windows , office .
cheap ...
4          2030   ham  Subject: re : indian springs\r\nthis deal is
t...
...          ...    ...
...
5166         1518   ham  Subject: put the 10 on the ft\r\nthe
transport...
5167          404   ham  Subject: 3 / 4 / 2000 and following noms\r\
nhp...
5168         2933   ham  Subject: calpine daily gas nomination\r\n>\r\
n...
5169         1409   ham  Subject: industrial worksheets for august
2000...
5170         4807  spam  Subject: important online banking alert\r\
ndea...
```

```
      label_num
0              0
1              0
2              0
3              1
4              0
...          ...
5166           0
5167           0
5168           0
5169           0
5170           1
```

```
[5171 rows x 4 columns]>
```

```
x=plt.xlabel("label")
y=plt.ylabel("count")
s.plot.bar()
```

```
<Axes: xlabel='label', ylabel='count'>
```



```
data=data.drop('Unnamed: 0',axis=1)
data=data.drop('label',axis=1)
```

```
data.head()
```

	text	label_num
0	Subject: enron methanol ; meter # : 988291\r\n...	0
1	Subject: hpl nom for january 9 , 2001\r\n( see...	0
2	Subject: neon retreat\r\nho ho ho , we ' re ar...	0
3	Subject: photoshop , windows , office . cheap ...	1
4	Subject: re : indian springs\r\nthis deal is t...	0

```
x=data['text']
```

```
y=data['label_num']
```

```
print(x.head())
```

```
0    Subject: enron methanol ; meter # : 988291\r\n...
1    Subject: hpl nom for january 9 , 2001\r\n( see...
2    Subject: neon retreat\r\nho ho ho , we ' re ar...
3    Subject: photoshop , windows , office . cheap ...
4    Subject: re : indian springs\r\nthis deal is t...
Name: text, dtype: object
```

```

from sklearn.feature_extraction.text import TfidfVectorizer

feature_extraction= TfidfVectorizer(min_df = 1 , stop_words='english',
lowercase = "True")

x=feature_extraction.fit_transform(x)

print(x)

(0, 37277)    0.2011097309413472
(0, 17826)    0.2392098021343849
(0, 7216)     0.1904342633768622
(0, 32458)    0.15611810136067314
(0, 11853)    0.12297449106504706
(0, 13882)    0.16097016125305302
(0, 21604)    0.09031870628106299
(0, 33380)    0.22464796050060742
(0, 5141)     0.16330305397181255
(0, 38455)    0.18654788930756974
(0, 49869)    0.17077874338809632
(0, 36555)    0.2011097309413472
(0, 47605)    0.11813966930318913
(0, 15070)    0.26506114668840103
(0, 36121)    0.2304064936785402
(0, 34299)    0.27312912121289434
(0, 15168)    0.09383338654057037
(0, 37033)    0.17198604767379225
(0, 15225)    0.16554063068104793
(0, 20643)    0.12642883285390433
(0, 36490)    0.20964545173769655
(0, 0)        0.10426983213156465
(0, 31571)    0.14947169013950107
(0, 21659)    0.20324385893640598
(0, 33042)    0.14755582683265941
:            :
(5170, 40944) 0.043825552681914126
(5170, 19502) 0.06940047912119972
(5170, 48113) 0.06743231879609364
(5170, 5896)  0.07807203338816637
(5170, 39943) 0.0654536507459819
(5170, 25582) 0.04084595884272012
(5170, 25931) 0.060842568812869834
(5170, 41025) 0.11074677663035766
(5170, 7528)  0.06905149192063627
(5170, 268)   0.034142006565318614
(5170, 21370) 0.05574695522298963
(5170, 6815)  0.07327442668249698
(5170, 12642) 0.05049708257577212
(5170, 29821) 0.045788574177944755

```

```
(5170, 33780) 0.32156994184710264
(5170, 32467) 0.0520749217466965
(5170, 30849) 0.13139963528249582
(5170, 13830) 0.060033187232779735
(5170, 31235) 0.07559776400175358
(5170, 25157) 0.06481291025635455
(5170, 20756) 0.0445653878572119
(5170, 16946) 0.05332923495537823
(5170, 44941) 0.04048151839656568
(5170, 33042) 0.05448307092167786
(5170, 43337) 0.013260748263266092
```

```
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2)
```

```
print(x_train)
```

```
(0, 42103) 0.14380214994766397
(0, 41181) 0.14380214994766397
(0, 34615) 0.14380214994766397
(0, 31965) 0.15070039628782436
(0, 48718) 0.15070039628782436
(0, 45874) 0.15070039628782436
(0, 2872) 0.15070039628782436
(0, 1361) 0.12331875735725041
(0, 3715) 0.15070039628782436
(0, 27641) 0.1293869283970034
(0, 13691) 0.15070039628782436
(0, 25581) 0.15070039628782436
(0, 18822) 0.15070039628782436
(0, 44106) 0.14380214994766397
(0, 22320) 0.15070039628782436
(0, 32404) 0.13511138736627842
(0, 12239) 0.14380214994766397
(0, 17221) 0.13890776627879634
(0, 8750) 0.13511138736627842
(0, 33122) 0.13890776627879634
(0, 42316) 0.13890776627879634
(0, 3509) 0.13890776627879634
(0, 13510) 0.13890776627879634
(0, 17536) 0.13200951993863597
(0, 15759) 0.13890776627879634
:
:
(4135, 43623) 0.046455556362546876
(4135, 18659) 0.05891424324849027
(4135, 46691) 0.06143311178984603
(4135, 1124) 0.05499698670773922
(4135, 2392) 0.04876028065672649
(4135, 38525) 0.060946929450399444
(4135, 33184) 0.08479234025357928
```

```
(4135, 46703) 0.050817292086376836
(4135, 25582) 0.03568628215580863
(4135, 41025) 0.04837860134334485
(4135, 7528) 0.06032888182274984
(4135, 40773) 0.06143311178984603
(4135, 29869) 0.04037489345974444
(4135, 11164) 0.1293125539276832
(4135, 16063) 0.051428993153682025
(4135, 12642) 0.0441182723554878
(4135, 14867) 0.05207480001491583
(4135, 46683) 0.04281988371384117
(4135, 24452) 0.036331632097748354
(4135, 11983) 0.04654678238993937
(4135, 48017) 0.04816601472532123
(4135, 36748) 0.09819550100861235
(4135, 44941) 0.03536787796212732
(4135, 33042) 0.09520149834712259
(4135, 43337) 0.023171291238002368
```

```
print(x_test)
```

```
(0, 19923) 0.09475824995592175
(0, 20446) 0.2007458914554433
(0, 29266) 0.1590903367294533
(0, 34715) 0.2246575631713854
(0, 12543) 0.0900989359816569
(0, 30431) 0.0762717644486883
(0, 24479) 0.0765168601146664
(0, 45222) 0.07247148825629958
(0, 34856) 0.0775430734918164
(0, 46022) 0.13359076841631123
(0, 34858) 0.08263515842555842
(0, 26091) 0.06890958341304583
(0, 34857) 0.06848779820870428
(0, 42842) 0.07157778826953606
(0, 29263) 0.07533258808731777
(0, 15393) 0.04888170976645469
(0, 34779) 0.07057632475352722
(0, 23464) 0.2683642441075446
(0, 22734) 0.05607918019934281
(0, 30125) 0.05788794887330025
(0, 40956) 0.09223354424495363
(0, 34718) 0.07382827623518655
(0, 47816) 0.0742416664446531
(0, 33452) 0.05868324167832812
(0, 40974) 0.3081873521007504
:
:
(1033, 2539) 0.12962188365420685
(1033, 12172) 0.2420459747205963
(1033, 6265) 0.2328720262376761
(1033, 2156) 0.1126084161472522
```

```
(1033, 198)    0.19374790249426826
(1033, 168)    0.20885915939156152
(1033, 14086)  0.08045495712583124
(1033, 2795)   0.13335350241238764
(1033, 27151)  0.4482200770864063
(1033, 1112)   0.12320442982836503
(1033, 11517)  0.06675141128367587
(1033, 1526)   0.2121225559504335
(1033, 20943)  0.07650311950543334
(1033, 13093)  0.14516807071045587
(1033, 49235)  0.18591982788450814
(1033, 20288)  0.09077846287105402
(1033, 7500)   0.07721851691029148
(1033, 24360)  0.14245892398690754
(1033, 18639)  0.12646908080606306
(1033, 43337)  0.05589125800063714
(1034, 12460)  0.6324096984255231
(1034, 41596)  0.6603893669690137
(1034, 45132)  0.2886410839898137
(1034, 36624)  0.2715055614815425
(1034, 43337)  0.08315595017424587
```

```
print(y_test)
```

```
2713    0
63       0
2678    0
607     0
3990    0
..
1358    0
3668    0
4112    1
4747    0
546     1
```

```
Name: label_num, Length: 1035, dtype: int64
```

```
from sklearn.linear_model import LogisticRegression
```

```
model=LogisticRegression()
```

```
model.fit(x_train,y_train)
```

```
LogisticRegression(C=1.0 ,class_weight=None ,dual=False,fit_intercept=
True,intercept_scaling=1,
l1_ratio=None ,max_iter=100,multi_class='auto',n_jobs=None ,penalty='l
2',random_state=None, solver='lbfgs' , tol=0.0001, verbose=0,
warm_start=False)
```

```
LogisticRegression(penalty='l2')
```

```
y_predict=model.predict(x_test)
```

```
print(y_predict)
```

```
[0 0 0 ... 1 0 1]
```

```
from sklearn.metrics import accuracy_score  
print(accuracy_score(y_test,y_predict)*100,'%')
```

```
98.93719806763285 %
```

```
# make prediction
```

```
input_mail=["Hi Yakesh K,World's biggest coding contest is backAre you  
geared up and excited to participatein this event?YesMay be laterAre  
you ready to compete against millions of developers and set new  
records? Join TechGig's flagship event, Code Gladiators, the biggest  
online tech community in India."]
```

```
input_mail_feature=feature_extraction.transform(input_mail)
```

```
predict=model.predict(input_mail_feature)
```

```
print(predict)
```

```
if(predict[0]==1):
```

```
    print('Ham mail')
```

```
else:
```

```
    print("spam mail")
```

```
[1]
```

```
Ham mail
```